

Binary gates: fair three-model comparison on the balanced 100-comment validation set

Gate / model	Precision	Recall	Specificity	Balanced acc.	F1	MCC
Substantive / 5 mini	0.755	0.973	0.148	0.560	0.850	0.226
Substantive / 5.6	0.897	0.712	0.778	0.745	0.794	0.441
Substantive / 5.1	0.904	0.644	0.815	0.729	0.752	0.408
Explicit / 5 mini	0.870	0.833	0.679	0.756	0.851	0.497
Explicit / 5.6	0.909	0.694	0.821	0.758	0.787	0.466
Explicit / 5.1	0.933	0.583	0.893	0.738	0.718	0.430

N = 100 (25 comments per industry). All models are evaluated on identical comments; the 23 added healthcare comments were blind to human labels and notes.